

Probability distributions EBP038A05: 2020-2021

Assignments

Nicky D. van Foreest and student assistants

February 11, 2021

CONTENTS

1	Assignment 1	4	
1.1	Have you read well?	4	
1.2	Exercise at about exam level	5	
1.3	Coding skills	6	
1.4	Challenges, optional	9	
2	Assignment 2	10	
2.1	Have you read well?	10	
2.2	Exercises at about exam level	11	
2.3	Coding skills	12	
2.4	Challenges	15	
3	Hints	17	

GENERAL INFORMATION

Here we just provide the exercises of the assignments. For information with respect to grading we refer to the course manual.

Each assignment contains several sections. The first section is meant to help you read the book well and become familiar with definitions and concepts of probability theory. These questions are mostly simple checks, not at exam level, but lower. The second section contains some exercises at about the exam level to get you started. Here you have to derive and explain a solution, in mathematical notation. Most of the selected exercises of the book are also at about (or just a bit above) exam level. The third section is about coding skills. We explain the rationale presently. The final section with challenges is for those students that like a challenge; the problems are above exam level.

You have to get used to programming and checking your work with computers, for instance by using simulation. The coding exercises address this skill. You should know that much of programming is ‘monkey see, monkey do’. This means that you take code of others, try to understand it, and then adapt it to your needs. For this reason we include the code to answer the question. The idea is that you copy the code, you run it and include the numerical results in your report. You should be able to explain how the code works. For this reason we include questions in which you have explain how the most salient parts of the code works.

We include python and R code, and leave the choice to you what to use. In the exam we will also include both languages in the same problem, so you can stay with the language you like. You should know, however, that many of you will need to learn multiple languages later in life. For instance, when you have to access databases to obtain data about customers, patients, clients, suppliers, inventory, demand, lifetimes (whatever), you often have to use sql. Once you have the raw data, you process it with R or python to do statistics or make plots. (While I (=NvF) worked at a bank, I used Fortran for numerical work, AWK for string parsing and making tables, excel, SAS to access the database, and matlab for other numerical work, all next to each other. I got tired of this, so I went to using python as it did all of this stuff, but then within one language.) For your interest, based on the statistics [here](#) or [here](#), python scores (much) higher than R in popularity; if you opt for a business career, the probability you have to use python is simply higher than to have to use R.

You should become familiar with look up documentation on coding on the web, no matter your programming language of choice. Invest time in understanding the, at times, rather technical and terse, explanations. Once you are used to it, the core documentation is faster to read, i.e., less clutter. In the long run, it pays off.

The rules:

1. For each assignment you have to turn in a pdf document typeset in \LaTeX . Include a title, group number, student names and ids, and date.
2. We expect brief answers, just a sentence or so, or a number plus some short explanation. The idea of the assignment is to help you studying, not to turn you in a writer.
3. When you have to turn in a graph, provide decent labels and a legend, ensure the axes have labels too.

1 ASSIGNMENT 1

1.1 *Have you read well?*

Ex 1.1. In your own words, explain what is

1. a joint PMF, PDF, CDF;
2. a conditional PMF, PDF, CDF;
3. a marginal PMF, PDF, CDF.

Ex 1.2. We have two r.v.s $X, Y \in [0, 1]^2$ with the joint PDF $f_{X,Y}(x, y) = 2I_{x \leq y}$.

1. Are X and Y independent?
2. Compute $F_{X,Y}(x, y)$.

Ex 1.3. Correct (that is, is the following claim correct)? We have two continuous r.v.s X, Y . Even though the joint CDF factors into the product of the marginals, i.e., $F_{X,Y}(x, y) = F_X(x)F_Y(y)$, it is still possible in general that the joint PDF does not factor into a product of marginals PDFs of X and Y , i.e., $f_{X,Y}(x, y) \neq f_X(x)f_Y(y)$.

Ex 1.4. Consider $F_{X,Y}(x, y)/F_X(x)$. Write this expression as a conditional probability. Is this equal to the conditional CDF of X and Y ?

Ex 1.5. Let X be uniformly distributed on the set $\{0, 1, 2\}$ and let $Y \sim \text{Bern}(1/4)$; X and Y are independent.

1. Present a contingency table for the X and Y .
2. What is the interpretation of the column sums the table?
3. What is the interpretation of the row sums of the table?
4. Suppose you change some of the entries in the table. Are X and Y still independent?

Ex 1.6. Apply the chicken-egg story. A machine makes items on a day. Some items, independent of the other items, are failed (i.e., do not meet the quality requirements). What is N , what is p , what are the ‘eggs’ in this context, and what is the meaning of ‘hatching’? What type of ‘hatching’ do we have here?

Ex 1.7. Correct? We have two r.v.s X and Y on \mathbb{R}^+ . It is given that $F_{X,Y}(x, y) = F_X(x)F_Y(y)$ for $x, y \leq 1/3$. Then X and Y are necessarily independent.

Ex 1.8. I select a random guy from the street, his height $X \sim \text{Norm}(1.8, 0.1)$, and I select a random woman from the street, her height is $Y \sim \text{Norm}(1.7, 0.08)$. I claim that since I selected the man and the woman independently, their heights are independent. Briefly comment on this claim.

Ex 1.9. Correct? For any two r.v.s X and Y on \mathbb{R}^+ with marginals F_X and F_Y , it holds that $P\{X \leq x, Y \leq y\} = F_X(x)F_Y(y)$.

Ex 1.10. Theorem 7.1.11. What is the meaning of the notation $X|N = n$?

Ex 1.11. Correct? X, Y are two discrete r.v.s with CDF $F_{X,Y}$. We can compute the PDF as $\partial_x \partial_y F_{X,Y}(x, y)$.

1.2 Exercise at about exam level

Ex 1.12. This is about the simplest model for an insurance company that I can think of. We start with an initial capital $I_0 = 2$. The company receives claims and contributions every period, a week say. In the i th period, we receive a contribution X_i uniform on the set $\{1, 2, \dots, 10\}$ and a claim C_i uniform on $\{0, 1, \dots, 8\}$.

1. What is the interpretation of $\bar{I}_n = \min I_i : 0 \leq i \leq n$?
2. What is the meaning of $I_1 = I_0 + X_1 - C_1$?
3. What is the meaning of $I_2 = I_1 + X_2 - C_2$?
4. What is the interpretation of $I'_1 = \max I_0 - C_1, 0 + X_1$?
5. What is the interpretation of $I'_2 = \max I'_1 - C_2, 0 + X_2$?
6. What is $P\{I_1 < 0\}$?
7. What is $P\{I'_1 < 0\}$?
8. What is $P\{I_2 < 0\}$?
9. What is $P\{I'_2 < 0\}$?
10. Provide an interpretation in terms of the inventory of rice, say, at a supermarket for I_1 and I'_1 .

1.3 Coding skills

Ex 1.13. Use simulation to estimate the answer of BH.7.1. Run the code below and explain line 9 of python code or line 7 of the R code.

Then run the code for a larger sample, e.g, num=1000 or so, but remove the prints of a, b, and succes, because that will fill your screen with numbers you don't need. Only for small simulations such output is handy so that you can check the code.

Compare the value of the simulation to the exact value.

```

1 import numpy as np
2
3 np.random.seed(3)
4
5 num = 10
6
7 a = np.random.uniform(size=num)
8 b = np.random.uniform(size=num)
9 success = np.abs(a - b) < 0.25
10 print(a)
11 print(b)
12 print(success)
13 print(success.mean(), success.var())

```

```

1 set.seed(3)
2
3 num <- 10
4
5 a <- runif(num)
6 b <- runif(num)
7 success <- abs(a-b) < 0.25
8 a
9 b
10 success
11 paste(mean(success), var(success))

```

Challenge (not obligatory): If you like, you can include a plot of the region (in time) in which Alice and Bob meet, and put marks on the points of the simulation that were 'successful'.

Ex 1.14. Let $X \sim \text{Exp}(3)$. Find a simple expression for $P\{1 < X \leq 4\}$ and compute the value. Then use simulation to check this value. Finally, use numerical integration to compute this value. What are the numbers? Explain lines 11, 21 and 26 of the python code or lines 7, 17 and 23 of the R code.

```

1 import numpy as np
2 from scipy.stats import expon

```

```

3  from scipy.integrate import quad
4
5  labda = 3
6
7  X = expon(scale = 1 / labda).rvs(1000)
8  # print(X)
9  print(X.mean())
10
11 success = (X > 1) * (X < 4)
12 # print(success)
13 print(success.mean(), success.std())
14
15
16 def F(x): # CDF
17     return 1 - np.exp(-labda * x)
18
19
20 def f(x): # density
21     return labda * np.exp(-labda * x)
22
23
24 print(F(4) - F(1))
25
26 I = quad(f, 1, 4)
27 print(I)

```

```

1  labda <- 3
2
3  X <- rexp(1000, rate = labda)
4  # X
5  mean(X)
6
7  success <- (X > 1) * (X < 4)
8  # print(success)
9  paste(mean(success), sd(success))
10
11
12 CDF <- function(x) { # CDF
13     return(1 - exp(-labda * x))
14 }
15
16 f <- function(x) { # density
17     return(labda * exp(-labda * x))
18 }
19
20

```

```
21 CDF(4) - CDF(1)
22
23 I = integrate(f, 1, 4)
24 I
```

1.4 Challenges, optional

You are free to choose one of these problems, but of course you can do both if you like.

A UNIQUENESS PROPERTY OF THE POISSON DISTRIBUTION Consider again the chicken-egg story (BH 7.1.9): A chicken lays a random number of eggs N and each egg independently hatches with probability p and fails to hatch with probability $q = 1 - p$. Formally, $X|N \sim \text{Bin}(N, p)$. Assume also that $X|N \sim \text{Bin}(N, p)$ and that $N - X$ is independent of X . For $N \sim \text{Pois}(\lambda)$ it is shown in BH 7.1.9 that X and Y are independent. This exercise asks for the converse: showing that the independence of X and Y implies that $N \sim \text{Pois}(\lambda)$ for some λ . Hence, the Poisson distribution is quite special: it is the only distribution for which the number of hatched eggs doesn't tell you anything about the number of unhatched eggs.

Let $0 < p < 1$. Let N be an r.v. taking non-negative integer values with $P(N > 0) > 0$. Assume also that $X|N \sim \text{Bin}(N, p)$ and that $N - X$ is independent of X .

Ex 1.15. Use the assumption that $P\{N > 0\} > 0$ to prove that N has support \mathbb{N} , i.e. $P\{N = n\} > 0$ for all $n \in \mathbb{N}$. Note: $0 \in \mathbb{N}$.

Ex 1.16. Write $Y = N - X$. Prove that

$$P\{X = x\} P\{Y = y\} = \binom{x+y}{x} p^x (1-p)^y P\{N = x+y\}. \quad (1.1)$$

Ex 1.17. Prove that N is Poisson distributed.

IMPROPER INTEGRALS AND THE CAUCHY DISTRIBUTION This problem challenges your integration skills and lets you think about the subtleties of integrating a function over an infinite domain. (Such integrals are called improper integrals.)

Assume that X has the Cauchy distribution. Recall that $E[X]$ does not exist (hence, it is not automatic that the expectation of a some arbitrary r.v. exists).

Ex 1.18. Why does $E\left[\frac{|X|}{X^2+1}\right]$ exist? Find its value.

Ex 1.19. Explain why the previous exercise implies that $E\left[\frac{X}{X^2+1}\right]$ exists. Then find its value.

2 ASSIGNMENT 2

2.1 *Have you read well?*

Ex 2.1. Example 7.2.2. Write down the integral to compute $E[(X - Y)^2]$. You don't have to solve the integral.

Ex 2.2. Give a brief example of a situation where it might be more convenient to employ the correlation instead of the covariance and explain why.

Ex 2.3. In queueing theory the concept of squared coefficient of variance SCV of a rv X is very important. It is defined as $C = V[X]/(E[X])^2$. Is the SCV of X equal to $\text{Corr}(X, X)$? Can it happen that $C > 1$?

Ex 2.4. Using the definition of Covariance (Definition 7.3.1) derive the expression $\text{Cov}[X, Y] = E[XY] - E[X]E[Y]$. Use this finding to show why independence of X and Y implies their uncorrelatedness (note that the converse does not hold).

Ex 2.5. Let X, Y be two r.v.s and let a, b be two arbitrary scalars. Express $\text{Cov}[a(U + V), b(U - V)]$ in terms of $V[U]$, $V[V]$ and $\text{Cov}[U, V]$ (by using the expression obtained in the previous question).

Ex 2.6. Prove the key properties of covariance 1 to 5 on page 327 of the book (page 338 pdf).

Ex 2.7. Come up with a short illustrative example in which the random vector $\mathbf{X} = (X_1, \dots, X_6)$ follows a Multinomial Distribution with parameters $n = 10$ and $\mathbf{p} = (\frac{1}{6}, \dots, \frac{1}{6}) \in \mathbb{R}^6$.

Ex 2.8. Is the following claim correct? If the r.v.s X, Y are both normally distributed, then (X, Y) follows a Bivariate Normal distribution.

Ex 2.9. Let (X, Y) follow a Bivariate Normal distribution, with X and Y marginally following $\mathcal{N}(\mu, \sigma^2)$ and with correlation ρ between X and Y .

1. Use the definition of a Multivariate Normal Distribution to show that $(X + Y, X - Y)$ is also Bivariate Normal.
2. Find the marginal distributions of $X + Y$ and $X - Y$.
3. Compute $\text{Cov}[X, Y]$. Then, write down the expression for the joint PDF of $(X + Y, X - Y)$.

Ex 2.10. Let X, Y, Z be i.i.d. $\mathcal{N}(0, 1)$. Determine whether or not the random vector

$$\mathbf{W} = (X + 2Y, 3X + 4Z, 5Y + 6Z, 2X - 4Y + Z, X - 9Z, 12X + \sqrt{3}Y - \pi Z)$$

is Multivariate Normal. (Explain in words, don't do a lot of tedious math here!)

2.2 Exercises at about exam level

Ex 2.11. Take $X \sim \text{Unif}(\{-2, -1, 1, 2\})$ and $Y = X^2$. What is the correlation coefficient of X and Y ? If we would consider another distribution for X , would that change the correlation?

Ex 2.12. We have a machine that consists of two components. The machine works as long as not both components have failed. Let X_i be the lifetime of component i .

1. What is the interpretation of $\min\{X_1, X_2\}$?
2. If $X_1, X_2 \text{ iid} \sim \text{Exp}(10)$ (in hours), what is the probability that the machine is still 'up' (i.e., not failed) at time T ?
3. Use the previous result to determine the distribution of $\min\{X_1, X_2\}$.
4. What is the expected time until the machine fails?

Ex 2.13. We have two r.v.s X and Y with the joint PDF $f_{X,Y}(x,y) = \frac{6}{7}(x+y)^2$ for $x, y \in (0, 1)$ and 0 else. Also we consider the two r.v.s U and V with the joint PDF $f_{U,V}(u,v) = 2$ for $u, v \in [0, 1], u+v \leq 1$ and 0 else.

1. Compute $P\{X+Y > 1\}$.
2. Compute $\text{Cov}[U, V]$.

(Hint: first draw the area over which you want to integrate, if this does not help check out the discussion board post on exercise 7.13a from the first Tutorial)

2.3 Coding skills

VERIFY THE ANSWERS OF BH.5.6.5 Read this example of BH first. We chop up the exercise in many small exercises..

For python code below, run it for a small number of samples; here I choose samples=2. Read the print statements, and use that to answer the questions below.

```

1 import numpy as np
2 from scipy.stats import expon
3
4 np.random.seed(10)
5
6 labda = 4
7 num = 3
8 samples = 2
9
10 X = expon(scale=labda).rvs((samples, num))
11 print(X)
12 T = np.sort(X, axis=1)
13 print(T)
14 print(T.mean(axis=0))
15
16 expected = np.array([labda / ((num - j)) for j in range(num)])
17 print(expected)
18 print(expected.cumsum())

```

Ex 2.14. In line P.11¹ we print the value of X in line 10. What is the meaning of X?

Ex 2.15. What is the meaning of T in line P.12?

Ex 2.16. What do we print in line P.14?

Ex 2.17. What is meaning of the variable expected?

Ex 2.18. What is the cumsum of expected?

Ex 2.19. Now that you understand what is going on, rerun the simulation for a larger number of samples, e.g., 1000, and discuss the results briefly.

ON BH.7.48 Read this exercise first and solve it. Then consider the code below.

```

1 import numpy as np
2
3 np.random.seed(3)
4
5

```

¹ Line P.x refers to line x of the Python code. Line R.x refers to line x of the R code.

```

6 def find_number_of_maxima(X):
7     num_max = 0
8     M = -np.infty
9     for x in X:
10         if x > M:
11             num_max += 1
12             M = x
13     return num_max
14
15
16 num = 10
17 X = np.random.uniform(size=num)
18 print(X)
19
20 print(find_number_of_maxima(X))
21
22 samples = 100
23 Y = np.zeros(samples)
24 for i in range(samples):
25     X = np.random.uniform(size=num)
26     Y[i] = find_number_of_maxima(X)
27
28 print(Y.mean(), Y.var(), Y.std())

```

Ex 2.20. Explain how the small function in lines P.6 to P.13 works. (You should know that $x += 1$ is an extremely useful abbreviation of the code $x = x + 1$).

Ex 2.21. Explain the code in lines P.25 and P.26.

WHY IS THE EXPONENTIAL DISTRIBUTION SO IMPORTANT? At the Paris metro, a train arrives every 3 minutes on a platform. Suppose that 50 people arrive between the departure of a train and an arrival. It seems entirely reasonable to me to model the arrival times of the individual people as distributed on the interval $[0, 3]$. What is the distribution of the inter-arrival times of these people? It turns out to be exponential! Story BH.13.4.2 provides an explanation (It is not forbidden to read the book beyond what you have to do for this course!); here we use simulation to check this fact.

```

1 import numpy as np
2
3 np.random.seed(3)
4
5
6 num = 5 # small sample at first, for checking.
7 start, end = 0, 3
8 labda = num / (end - start) # per minute
9 print(1 / labda)
10

```

```

11 A = np.sort(np.random.uniform(start, end, size=num))
12 print(A)
13 print(A[1:])
14 print(A[:-1])
15 X = A[1:] - A[:-1]
16 print(X)
17
18 print(X.mean(), X.std())

```

Ex 2.22. Explain the result of line P.12.

Ex 2.23. Compare the result of line P.12 and P.13; explain what is $A[1:]$.

Ex 2.24. Compare the result of line P.12 and P.14; explain what is $A[:-1]$.

Ex 2.25. Explain what is X in P.16.

Ex 2.26. Why do we compare $1/\lambda$ and $X.\text{mean}()$?

Ex 2.27. Recall that $E[X] = \sigma(X)$ when $X \sim \text{Exp}(\lambda)$. Hence, what do you expect to see for $X.\text{std}()$?

Ex 2.28. Run the code for a larger sample, e.g. 50, and discuss (very briefly) your results.

2.4 Challenges

This exercise will give an example of how probability theory can pop up in OR problems, in particular in linear programs. It introduces you to the concept of *recourse models*, which you will learn about in the master course Optimization Under Uncertainty. **Disclaimer: the story is quite lengthy, but the concepts introduced and questions asked are in fact not very hard. We just added the story to make things more intuitive.**

We consider a pastry shop that only sells one product: chocolate muffins. Every morning at 5:00 a.m., the shop owner bakes a stock of fresh muffins, which he sells during the rest of the day. Making one muffin comes at a cost of $c = \$1$ per unit. Any leftover muffins must be discarded at the end of the day, so every morning he starts with an empty stock of muffins.

The owner has one question for you: determine the amount x of muffins that he should make in the morning to minimize his production cost. Note that the owner never wants to disappoint any customer, i.e., he requires that $x \geq d$, where d is the daily demand for muffins.

The problem can be formulated as a linear program (LP):

$$\min_{x \geq 0} \{cx : x \geq d\}. \quad (2.1)$$

For simplicity, we ignore the fact that x should be integer-valued.

Ex 2.29. Determine the optimal value x^* for x and the corresponding objective value in case d is deterministic.

Of course, in practice d is not deterministic. Instead, d is a random variable with some distribution. However, note that the LP above is ill-defined if d is a random variable. We cannot guarantee that $x \geq d$ if we do not know the value of d .

You explained the issue to the shop owner and he replies: “Of course, you’re right! You know, whenever I’ve run out of muffins and a customer asks for one, I make one on the spot. I never disappoint a customer, you know! It does cost me 50% more money to produce them on the spot, though, you know.”

Mathematically speaking, the shop owner just gave you all the (mathematical) ingredients to build a so-called *recourse model*. We introduce a *recourse variable* y in our model, representing the amount of muffins produced on the spot. Production comes at a unit cost of $q = 1.5c = \$1.5$. Assuming that we know the distribution of d , we can then minimize the *expected total cost*:

$$\min_{x \geq 0} \{cx + E[v(d, x)]\}, \quad (2.2)$$

where $v(d, x)$ is the optimal value of another LP, namely the *recourse problem*:

$$v(d, x) := \min_{y \geq 0} \{qy : x + y \geq d\}, \quad (2.3)$$

for given values of d and x . The recourse problem can easily be solved explicitly: we get $y = d - x$ if $d \geq x$ and $y = 0$ if $d < x$. So we obtain

$$v(d, x) = q(d - x)^+, \quad (2.4)$$

where the operator $(\cdot)^+$ represents the *positive value operator*, defined as

$$(s)^+ = \begin{cases} s & \text{if } s \geq 0, \\ 0 & \text{if } s < 0. \end{cases} \quad (2.5)$$

Ex 2.30. To get some more insight into the model, suppose (for now) that $d \sim U\{10, 20\}$. Solve the model, i.e., find the optimal amount x^* . *Hint: First, compute the value of $E[v(d, x)]$ as a function of x . Then find the optimal value of x .*

Ex 2.31. What is the expected recourse cost (expected cost of on-the-spot production) at the optimal solution x^* , i.e., compute $E[v(d, x^*)]$?

To solve the model correctly, we need the true distribution of d . We learn the following from the shop owner: “My granddaughter, who’s always running around in my shop, is a bit data-crazy, you know, so she’s been collecting some data. I remember her saying that ‘the demand from male and female customers are both approximately normally distributed, with mean values both equal to 10 and standard deviations of 5’. She also mentioned something about correlation, but I don’t remember exactly, you know. It was either almost 1 or almost -1 . I hope this helps!”

Mathematically, we’ve learned that $d = d_m + d_f$, with $(d_m, d_f) \sim \mathcal{N}(\mu, \Sigma)$, where $\mu = (\mu_m, \mu_f) = (10, 10)$ and $\Sigma_{11} = \sigma_m^2 = \Sigma_{22} = \sigma_f^2 = 5^2 = 25$. Finally, $\Sigma_{12} = \Sigma_{21} = \text{Cov}(d_m, d_f) = \rho \sigma_m \sigma_f = 25\rho$. Also, we know that either $\rho \approx 1$ or $\rho \approx -1$.

Ex 2.32. Calculate x^* and the corresponding objective value for the case $\rho = -1$. (Do not read $\rho = 1$, this case is not simple.)

Ex 2.33. Consider the two extreme cases $\rho = 1$ and $\rho = -1$. In which case will the shop owner have lower expected total costs? Provide a short, intuitive explanation. *Hint: you don’t have to compute x^* for the case where $\rho = 1$ (this is not easy!).*

3 HINTS

h.1.8. From this exercise you should memorize this: **independence is a property of the joint CDF, not of the rvs.**

h.1.17. Use the relation of the previous exercise to show that

$$P(N = n + 1) = \frac{\lambda}{1 + n} P(N = n). \quad (3.1)$$

Bigger hint: Fill in $y = 0$ in the LHS and RHS of (1.1); call this expression 1. Then fill in $y = 1$ to obtain a second expression. Divide these two expressions and note that $P\{X = x\}$ cancels. Finally, define

$$\lambda = \frac{P\{Y = 1\}}{(1 - p)P\{Y = 0\}}. \quad (3.2)$$